Balancing consistency with interpretive depth

Reasoning models produced dependable answers, but often stopped at surface-level meaning. More creative models could identify context, tensions, and hidden signals, yet applied their logic less consistently across repeated datasets.

To evaluate the trade-off, the designer ran blind tests using the same instructions, reference documents, and source data, then compared both repeatability and interpretive depth. The goal stayed fixed. What changed were the prompts and supporting materials: definitions were clarified, examples were added, and reference documents were refined after each test.

The iterations showed that neither reliability nor creativity was sufficient alone. GPT-o3 established a strong consistency benchmark, while GPT-4.5 demonstrated the intuitive reading needed for strategic research. Combining o3-like structure with GPT-4.5's interpretive range moved the pipeline toward answers that were both steadier across repeated runs and better able to recognize underlying meaning.

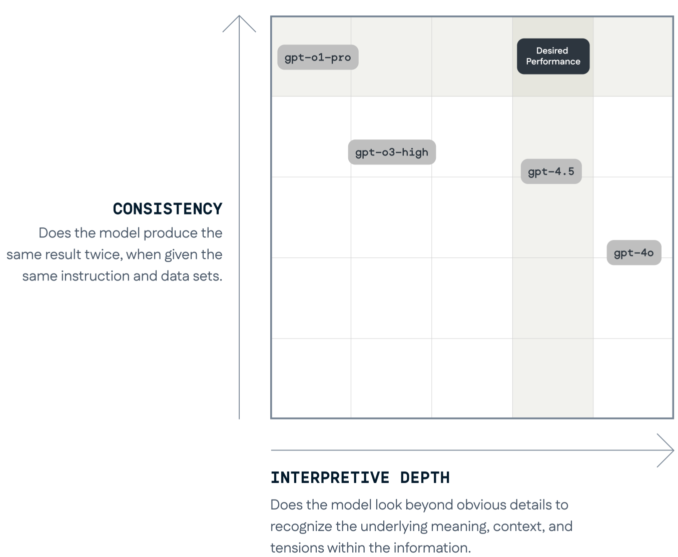

DEGREES OF INTERPRETIVE DEPTH

Refining the same goal across repeated tests

The table tracks how consistency and interpretive depth changed as the prompt and reference materials evolved.

Improving consistency and interpretive depth through prompting, measured using blind tests.

Iterations	gpt-4o		gpt-4.5		gpt-o3-high		gpt-o1-pro	
Iteration 01	~36%	Imaginative	~50%	Intuitive	~65%	Well-Reasoned	~95%	Literal
Iteration 02	~50% ↑	Imaginative	~62% ↑	Intuitive	~65%	Well-Reasoned	~95%	Literal
Iteration 03	~40% ↓	Imaginative	~54% ↓	Imaginative ↑	~75% ↑	Well-Reasoned	discontinued	
Iteration 04	~45% ↑	Imaginative	~72% ↑	Well-Reasoned ↓	~85% ↑	Well-Reasoned		
Iteration 05	discontinued		~84% ↑	Intuitive ↑	~85%	Well-Reasoned		
Iteration 06			~76% ↓	Well-Reasoned ↓	discontinued			
Iteration 07			~88% ↑	Well-Reasoned				
Iteration 08			~72% ↓	Well-Reasoned				
	Consistency Improved. But, the interpretation remained too abstract and imaginative.		Required multiple iterations to identify the perfect balance between consistency and interpretive depth.		While consistency could be improved, interpretations remained ‘Well-Reasoned’ but not ‘Intuitive’.		gpt-o1 helped establish a benchmark for consistency, but improving interpretive depth proved challenging.	

Iteration 05 produced the strongest combined result: GPT-4.5 reached approximately 84% consistency while remaining intuitive. Later declines showed that gains were not automatic; prompts and reference documents needed to operate as one deliberately designed system.