


PANW - Defining AI interaction patterns for spreadsheet-native users.




I led the end-to-end design of an AI supported sales-operations workbench by rethinking:

- How we design data-dense interactive tables.
- How AI converses with spreadsheet-native users.
- What “source” means for internal teams.

Problem #1

Teams had a scattered landscape of multiple single-purpose tools built by distinct teams. The typical approach to consolidation involved a 15+ column table.




Decisions

I redesigned the table experience by using interactive tags as the atomic unit of data instead of cells, and thematically grouping them.

To prevent monotony or excessive cognitive load, I kept the tags plain but added colorful iconography to signal distinct meanings.

Outcome


Post-launch feedback led to the cluster table becoming the default pattern for future table designs with multiple columns.




Problem #2


The project began with a push to bring AI into the sales workflow and specifically as a chatbot. The real disconnect was between the narrative form that AI takes and the fluency users had formed over decades of industry practice.

A FLUENCY SHAPED BY DECADES OF PRACTICE AND INDUSTRY NORMS.




THE INTERACTION PATTERN SEEN AS A NECESSITY FOR PROGRESS.





Decisions

By experimenting with CopilotKit, I designed tabular responses with interactive rows and columns to enable follow-up interactions. This allowed the users to engage in exploratory analysis rather than single-prompt interactions.



Outcome

Through observation, we found users using this combination as a means to triangulate root causes, which led me to the placeholder tagline:


"Ask what usually takes four tools to answer."

In later phases, we also expanded the interaction pattern to allow context from UI tables to trigger new chats.

Problem #3

Due to conventional CUI experiences, "thinking" and "sources" were treated as an afterthought. However 2 observations led me to rethink these critical features.

1. While the typical "thinking" shows a meandered path models take to reach a conclusion, users needed a structured "derivation" of a response, leading with the conclusion.
2. A platform's name does not pass as a "source" for internal teams. What's more important:
 - Which dataset within the platform?
 - When was the data last updated?




Decision

These observations led me to design Trace (renamed from 'Derivation')

Trace breaks down a model's thinking into 4–7 concrete steps, with pathways and timestamps for last refresh.

Outcome

In practice, we found that team members would snap a response and its trace together, and drop it in Slack threads to defend a number. The pattern became frequent enough that snapshotting was promoted to an enhancement.



Lessons

Progress lies somewhere between blind acceptance and complete dismissal.

It was certainly uncomfortable being yet another team trying to add a chatbot with tiny chips that simulate "thinking" and treat citations as checkboxes. But acknowledging discomfort while taking action led to some of the more unique experiences in the application, with a lasting impact.

The question was never "how do we add a chatbot?" but rather "what assumptions about the chatbot experience do we carry that will fail our users?"

